# MISSION: Phase 8 Feature Prioritization

**Mission ID:** MISSION-2024-PHASE8
**Date:** December 2, 2024
**Status:** Planning
**Priority:** High

---

## Mission Overview

**Title:** Determine Phase 8 Feature Roadmap using Virtual CEO Framework

**Mission Type:** Strategic Initiative

**Description:**
Analyze 10+ candidate features from Phase 7B+ backlog to determine the optimal Phase 8 roadmap. Use Virtual CEO multi-perspective analysis to prioritize features that maximize business value, technical feasibility, user impact, and competitive positioning.

---

## Context

### Business Context

- **Current State:** Phase 7A complete (Portfolio Analytics)
- **Revenue:** Targeting $100k MRR by Q2 2025
- **Current MRR:** ~$15k (estimated, based on subscriber base)
- **Growth Rate:** 15-20% MoM
- **User Base:** ~50 organizations across Collector, Dealer, Enterprise tiers
- **Competitive Pressure:** 2-3 direct competitors launching similar features

### Technical Context

- **Tech Stack:** Next.js, PostgreSQL, AWS S3, Stripe, Google Vision AI
- **Infrastructure:** Scalable (ECS, RDS)
- **Team Size:** 1-2 developers (+ AI assistance)
- **Tech Debt:** Low (clean codebase, 0 TypeScript errors)
- **Current Build:** 66 routes, production-ready

### Constraints

- **Budget:** $30k for Phase 8 (development + external services)
- **Timeline:** 8-12 weeks target
- **Resources:** 1-2 developers, part-time product management
- **Must maintain:** Zero downtime during development

---

# Candidate Features (Backlog)

## From Phase 7B+ Planning:

1. **Predictive Analytics (ML-powered value predictions)**
2. **Advanced Visualizations (heat maps, correlation charts)**
3. **Custom Report Templates**
4. **Benchmark Comparisons (market indices)**
5. **AI Insights (automated trend detection)**
6. **Push Notifications (PWA)**
7. **Offline Sync**
8. **VIN OCR (camera-based VIN reading)**
9. **Enhanced QR Codes**
10. **International VIN Support**
11. **Marketplace Integration (eBay, Chrono24, etc.)**
12. **Public Asset Profiles (social sharing)**
13. **Insurance Partner API**
14. **Auction House Integration**
15. **Mobile App (Native iOS/Android)**

---

# Success Criteria

## Primary Goals

1. **Feature Prioritization:** Top 3-5 features ranked with clear, data-driven rationale
2. **Detailed Execution Plan:** Phase 1 feature (#1 priority) with:
   - Week-by-week breakdown
   - Resource allocation
   - Technical design
   - Success metrics
3. **Risk Assessment:** Identified risks with mitigation strategies for top 3 features
4. **Resource Allocation:** Clear budget and timeline for each prioritized feature

## Success Metrics

- **Decision Speed:** Complete prioritization within 1 week
- **Stakeholder Alignment:** 100% team buy-in on #1 priority
- **Execution Confidence:** >80% confidence in estimates and approach
- **Business Value:** Selected features collectively target +$25k MRR impact

---

# Virtual CEO Multi-Perspective Analysis

## Scoring Framework

Each feature will be scored (1-10 scale) across 7 perspectives:

**CEO (Strategic Impact)** - 25% weight
- Competitive differentiation
- Market opportunity
- Mission alignment
- Long-term vision

**CTO (Technical Feasibility)** - 20% weight
- Complexity (inverse)
- Architecture fit
- Maintainability
- Scalability

**CPO (User Value)** - 20% weight
- User demand/requests
- Problem severity
- UX improvement
- Feature adoption likelihood

**CFO (Financial Impact)** - 15% weight
- Revenue opportunity
- Cost to build
- ROI timeline
- Operating costs

**COO (Operational Impact)** - 10% weight
- Support burden
- Infrastructure needs
- Vendor dependencies
- Deployment complexity

**CMO (Marketing Value)** - 5% weight
- Differentiation
- PR opportunity
- Customer acquisition
- Positioning

**CISO (Security/Compliance)** - 5% weight
- Security risks
- Compliance impact
- Data privacy
- Audit requirements

**Total Score:** Weighted sum (max 10.0)

---

# Feature Analysis Matrix

## Feature 1: Predictive Analytics (ML-Powered Value Predictions)

### CEO Perspective (Score: 9/10)

**Strategic Impact:**
- **Competitive Differentiation:** HIGH - No competitor offers ML predictions for luxury assets

- **Market Opportunity:** $50B luxury resale market hungry for data-driven insights
- **Mission Alignment:** Perfect fit - “AI-native platform for provenance”
- **Long-Term Vision:** Foundation for recommendation engine, portfolio optimization, fraud detection

**Analysis:**
This is a game-changer. Positions us as the “Bloomberg Terminal for luxury assets.” Dealers and collectors constantly ask “What will this be worth in 3 years?” - we’d be the only ones answering with AI. Creates massive moat (6-12 month lead on competitors). Aligns with Series A fundraising narrative.

**Concerns:** Accuracy is critical - predictions must be >80% accurate or it damages brand.

---

## CTO Perspective (Score: 6/10)

**Technical Feasibility:**
- **Complexity:** MEDIUM-HIGH (ML model training, external data integration, retraining pipeline)
- **Architecture Fit:** GOOD (new microservice, doesn’t impact existing platform)
- **Maintainability:** MEDIUM (requires ongoing model monitoring and retraining)
- **Scalability:** GOOD (predictions can be cached, async processing)

**Technical Requirements:**
1. ML Infrastructure:
- AWS SageMaker or Google Vertex AI
- Model training pipeline
- Automated retraining (monthly)
- A/B testing framework

1. Data Integration:
   - Chrono24 API (watches)
   - Hemmings/Bring a Trailer (cars)
   - Artsy/Artnet (art)
   - Historical price data storage

2. New Services:
   - Prediction API service
   - Model monitoring service
   - Data ingestion pipeline

**Estimated Effort:** 8-10 weeks
- Week 1-2: Data partnerships + ingestion pipeline
- Week 3-6: Model development + training
- Week 7-8: API + UI integration
- Week 9-10: Testing + monitoring setup

**Technical Risks:**
- Data quality from external APIs
- Model accuracy (target: 80%, may achieve 70-75% initially)
- API rate limits from data providers
- Cold start problem (need historical data)

**Mitigation:**
- Start with watches only (best data availability)
- Partner with Chrono24 directly

- Set clear accuracy thresholds before launch
- Gradual rollout to Dealer tier first

---

## CPO Perspective (Score: 9/10)

**User Value:**
- **User Demand:** HIGH - #1 requested feature from Dealer tier (18 requests in last 3 months)
- **Problem Severity:** HIGH - Dealers making $100k+ decisions with no data
- **UX Improvement:** SIGNIFICANT - Transforms analytics from reporting to actionable insights
- **Adoption Likelihood:** HIGH - Natural extension of existing analytics (47% use portfolio trends weekly)

**User Stories:**
1. **As a dealer**, I want predicted appreciation so I can optimize buying decisions
2. **As a collector**, I want predicted value so I can plan selling timing
3. **As an auction house**, I want predictions so I can set better reserves

**MVP Scope:**
- Predictions for watches, cars, art only
- 1-year, 3-year, 5-year predictions
- Confidence intervals (±15%)
- Historical accuracy tracking
- Simple UI: card in analytics dashboard

**Full Scope (Post-MVP):**
- All categories
- Market trend explanations (“Value increasing due to...”)
- Comparison to similar assets
- Alerts on significant prediction changes
- Recommendation engine (“Buy now” vs “Wait”)

**User Engagement Prediction:**
- 60%+ of Dealer users view predictions weekly
- 30%+ share predictions with clients
- 15%+ cite predictions in upgrade decision

---

## CFO Perspective (Score: 8/10)

**Financial Impact:**

**Revenue Opportunity:**
- **Collector → Dealer Upgrades:** +15 upgrades @ $599/year = +$8,985 ARR
- **New Dealer Signups:** +10 new @ $599/year = +$5,990 ARR
- **Enterprise Upsell:** Feature in Enterprise tier, drives 5 upgrades @ $1,999/year = +$9,995 ARR
- **Total New ARR:** $24,970 (~$2,080/month)
- **Year 1 Impact:** +$25k ARR (conservative), +$50k ARR (aggressive)

**Cost to Build:**
- Development: $18,000 (10 weeks @ $1,800/week blended rate)
- Data Partnerships: $2,000 (Chrono24 API setup)

- ML Infrastructure: $500/month AWS SageMaker
- Data APIs: $200/month (Chrono24 + others)
- **Total First Year Cost:** $18,000 + $2,000 + ($700 × 12) = $28,400

**ROI Analysis:**
- **Conservative Scenario:** $25k ARR - $28.4k cost = -$3.4k Year 1, +$16.6k Year 2 (ROI: 119% over 2 years)
- **Aggressive Scenario:** $50k ARR - $28.4k cost = +$21.6k Year 1 (ROI: 176% Year 1)
- **Payback Period:** 13 months (conservative), 7 months (aggressive)

**Operating Costs (Ongoing):**
- AWS SageMaker: $500/month
- Data APIs: $200/month
- Model retraining: $100/month (compute)
- **Total:** $800/month = $9,600/year

**Unit Economics:**
- Cost per prediction: ~$0.05 (amortized)
- Value per prediction: N/A (bundled in subscription)
- Incremental margin: 98% (software)

**Financial Risks:**
- Data partnership costs increase
- Lower than expected conversion
- Competitor launches similar feature (reduces differentiation)

---

## COO Perspective (Score: 7/10)

**Operational Impact:**

**Support Burden:** MEDIUM
- New support tier: “Prediction Accuracy” queries
- Training required for support team
- FAQ documentation needed
- Estimated: +10 support tickets/week

**Infrastructure Needs:** MEDIUM
- AWS SageMaker (managed service)
- Additional database tables (predictions, model metrics)
- Monitoring dashboard (model performance)
- No new servers required

**Vendor Dependencies:** MEDIUM
- Chrono24 API (primary data source)
- AWS SageMaker (ML infrastructure)
- Backup: Manual data collection if APIs fail

**Deployment Complexity:** LOW
- Microservice architecture (isolated)
- Feature flag for gradual rollout
- Dealer tier first, then Enterprise
- Can roll back without data loss

**Operational Risks:**
- Chrono24 API downtime (mitigation: cached predictions)
- Model drift (mitigation: automated monitoring + alerts)
- Prediction disputes (mitigation: clear disclaimers + confidence intervals)

**Resource Requirements:**
- DevOps: 5 hours (infrastructure setup)
- Support Training: 2 hours
- Documentation: 4 hours
- Ongoing Monitoring: 2 hours/week

---

## CMO Perspective (Score: 8/10)

**Marketing Value:**

**Differentiation:** HIGH
- **Unique Selling Proposition:** “Only provenance platform with AI-powered value predictions”
- **Competitive Moat:** 6-12 month lead (complex to replicate)
- **Industry First:** Position as innovator

**PR Opportunity:** HIGH
- Press release: “Genesis Provenance Launches AI Predictions”
- Tech blogs: AI + luxury crossover story
- Industry publications: Robb Report, Hodinkee, Rennlist
- Podcast interviews: Collector/dealer shows

**Customer Acquisition:** MEDIUM
- SEO: “luxury watch value prediction AI”
- Content marketing: “How to predict watch appreciation”
- Case studies: Dealer success stories
- Webinar: “AI-powered collecting”

**Positioning:**
- **Current:** “Digital provenance for luxury assets”
- **New:** “AI-powered platform for provenance + predictions”
- Aligns with fundraising narrative
- Sets stage for future AI features

**Marketing Campaign:**
1. **Pre-Launch (2 weeks):**
- Teaser campaign to existing users
- Email to Dealer tier: “Coming soon: AI predictions”
- Social media countdown

1. **Launch (Week 1):**
   - Press release + media outreach
   - Blog post: Technical deep-dive
   - Email campaign to all users
   - LinkedIn + Twitter announcements

2. **Post-Launch (4 weeks):**
   - User success stories

- Demo videos
- Partner co-marketing (Chrono24?)
- Paid ads targeting dealers

**Estimated Marketing Budget:** $5,000
- PR distribution: $1,000
- Content creation: $2,000
- Paid ads: $2,000

---

## CISO Perspective (Score: 9/10)

**Security & Compliance:**

**Security Risks:** LOW
- **Data Flow:** Read-only external APIs → ML model → Display to users
- **No PII:** Predictions based on asset data, not user data
- **No Payment Data:** Purely informational feature

**Compliance Considerations:**
- **Financial Advice Disclaimer:** Required - predictions are informational, not financial advice
- **GDPR:** N/A (no personal data)
- **CCPA:** N/A (no personal data)
- **SOC 2:** Model monitoring logs + audit trail (already planned)

**Model Security:**
- **Model Poisoning:** Low risk (external data from reputable sources)
- **Model Theft:** Medium risk (mitigate with API rate limiting)
- **Model Explainability:** Required for transparency

**Audit Requirements:**
- Log all predictions (asset ID, timestamp, prediction, confidence)
- Track model performance (accuracy over time)
- Document model training data sources
- Version control for models

**Data Privacy:**
- No user data sent to external APIs
- Asset data (brand, model, year) is non-sensitive
- Predictions stored in secure database

**Liability Mitigation:**
- Clear disclaimers: “Predictions are estimates, not guarantees”
- Display confidence intervals (±15%)
- Show historical accuracy metrics
- Terms of Service update (no liability for prediction accuracy)

**Security Recommendations:**
- API key rotation (Chrono24, AWS)
- Rate limiting on prediction API
- Monitoring for unusual prediction patterns
- Regular model audits

---

### Weighted Score Calculation: Predictive Analytics

| Perspective | Raw Score | Weight | Weighted Score |
|---|---|---|---|
| CEO | 9 | 25% | 2.25 |
| CTO | 6 | 20% | 1.20 |
| CPO | 9 | 20% | 1.80 |
| CFO | 8 | 15% | 1.20 |
| COO | 7 | 10% | 0.70 |
| CMO | 8 | 5% | 0.40 |
| CISO | 9 | 5% | 0.45 |
| **TOTAL** | | | **8.00** |

**Priority Rank:** #1 (Strong candidate)

---

## Abbreviated Analysis: Other Top Candidates

### Feature 2: VIN OCR (Camera-Based VIN Reading)

**CEO:** 6/10 - Good UX improvement, but not strategic differentiator
**CTO:** 8/10 - Technically straightforward (Google Vision OCR)
**CPO:** 7/10 - Nice-to-have, reduces friction
**CFO:** 5/10 - Low revenue impact, more of a feature than a product
**COO:** 9/10 - Easy to deploy, low support burden
**CMO:** 4/10 - Limited marketing value
**CISO:** 9/10 - Low security risk

**Weighted Score:** 6.55
**Priority Rank:** #5

---

### Feature 3: Marketplace Integration (eBay, Chrono24)

**CEO:** 8/10 - New distribution channel, high strategic value
**CTO:** 5/10 - Complex integrations, API limitations
**CPO:** 8/10 - High user demand from dealers
**CFO:** 7/10 - Potential revenue share model
**COO:** 5/10 - High support burden, vendor dependencies
**CMO:** 7/10 - Good partnership PR
**CISO:** 6/10 - API security concerns

**Weighted Score:** 6.90
**Priority Rank:** #3

## Feature 4: Push Notifications (PWA)

**CEO:** 5/10 - Table stakes, not differentiator
**CTO:** 7/10 - Moderate complexity (service workers)
**CPO:** 6/10 - Useful but not critical
**CFO:** 4/10 - Low revenue impact
**COO:** 6/10 - Notification management overhead
**CMO:** 3/10 - Minimal marketing value
**CISO:** 7/10 - Low security risk

**Weighted Score:** 5.55
**Priority Rank:** #7

---

## Feature 5: Insurance Partner API

**CEO:** 7/10 - B2B channel, high long-term value
**CTO:** 6/10 - Moderate complexity (API design)
**CPO:** 6/10 - Indirect user value
**CFO:** 8/10 - Enterprise revenue potential
**COO:** 7/10 - Moderate operational complexity
**CMO:** 8/10 - Strong partnership narrative
**CISO:** 6/10 - Data sharing compliance

**Weighted Score:** 6.95
**Priority Rank:** #2

---

## Feature 6: Mobile App (Native iOS/Android)

**CEO:** 6/10 - Market expectation, but PWA already works
**CTO:** 4/10 - High complexity, new codebase
**CPO:** 7/10 - Better UX than PWA
**CFO:** 5/10 - High cost ($50k+), uncertain ROI
**COO:** 4/10 - New deployment pipeline, app store overhead
**CMO:** 6/10 - App store visibility
**CISO:** 7/10 - Standard mobile security

**Weighted Score:** 5.65
**Priority Rank:** #6

---

## Final Feature Rankings

| Rank | Feature | Score | Rationale |
|---|---|---|---|
| 1 | **Predictive Analytics** | 8.00 | High strategic value, strong revenue impact, achievable |
| 2 | **Insurance Partner API** | 6.95 | B2B growth channel, enterprise revenue |
| 3 | **Marketplace Integration** | 6.90 | New distribution, user demand |
| 4 | **Auction House Integration** | 6.70 | Prestige partnerships, high-value users |
| 5 | **VIN OCR** | 6.55 | Quick win, improves UX |
| 6 | **Mobile App** | 5.65 | PWA sufficient for now |
| 7 | **Push Notifications** | 5.55 | Nice-to-have, not critical |
| 8 | **Public Asset Profiles** | 5.40 | Viral growth potential, but niche |
| 9 | **International VIN Support** | 5.25 | Geographic expansion enabler |
| 10 | **Enhanced QR Codes** | 4.90 | Incremental improvement |

---

## Recommended Phase 8 Roadmap

### Priority 1: Predictive Analytics (Weeks 1-10)

**Investment:** $28,400
**Expected Return:** +$25-50k ARR
**Risk:** Medium

**Execution Plan:** (See detailed breakdown below)

---

### Priority 2: Insurance Partner API (Weeks 11-16)

**Investment:** $15,000
**Expected Return:** +$30-60k ARR (Enterprise subscriptions)
**Risk:** Medium-High (sales cycle)

**Rationale:** After predictive analytics launches, we'll have a compelling story for insurance partners: "Verified provenance + predicted values = better underwriting."

---

### Priority 3: VIN OCR (Weeks 17-18)

**Investment:** $5,000
**Expected Return:** Improved conversion (5-10%)
**Risk:** Low

**Rationale:** Quick win to maintain momentum. Leverages existing Google Vision integration.

---

## Detailed Execution Plan: Predictive Analytics

### Phase 1: Data Foundation (Weeks 1-2)

**Week 1: Data Partnership Setup**
- **Tasks:**
1. Contact Chrono24 API sales
2. Negotiate data access agreement
3. Set up API credentials
4. Design data schema for historical prices
- **Deliverables:**
- Chrono24 API access
- Database schema: `asset_prices` table
- Data ingestion pipeline (basic)
- **Team:** 1 developer
- **Dependencies:** Chrono24 response time
- **Risk:** API costs higher than expected

**Week 2: Historical Data Collection**
- **Tasks:**
1. Build data ingestion service
2. Collect 6-12 months historical data
3. Clean and normalize data
4. Validate data quality
- **Deliverables:**
- 10k+ historical price points
- Data quality report
- Ingestion automation
- **Team:** 1 developer
- **Success Metric:** 90%+ data quality score

---

## Phase 2: ML Model Development (Weeks 3-6)

**Week 3: Feature Engineering**
- **Tasks:**
1. Define prediction features (brand, model, year, condition, provenance)
2. Create feature extraction pipeline
3. Build training dataset
- **Deliverables:**
- Feature engineering code
- Training dataset (80/20 split)
- **Team:** 1 developer + data scientist (consultant, $2k)

**Week 4-5: Model Training**
- **Tasks:**
1. Set up AWS SageMaker
2. Train baseline model (linear regression)
3. Train advanced models (XGBoost, neural network)
4. Compare model performance
- **Deliverables:**
- 3 trained models
- Accuracy comparison report
- Model selection decision
- **Team:** 1 developer + data scientist
- **Success Metric:** >75% accuracy (within ±15%)

**Week 6: Model Validation**
- **Tasks:**
1. Validate on held-out test set
2. Test edge cases
3. Document model limitations
4. Set up model versioning
- **Deliverables:**
- Validation report
- Model documentation
- Model registry
- **Go/No-Go Decision:** If accuracy <75%, iterate on features

---

## Phase 3: API & Integration (Weeks 7-8)

**Week 7: Prediction API Service**
- **Tasks:**
1. Build REST API for predictions
2. Implement caching layer
3. Add rate limiting
4. Write API documentation
- **Deliverables:**
- `/api/predictions/[itemId]` endpoint
- API documentation
- Postman collection
- **Team:** 1 developer

**Week 8: UI Integration**
- **Tasks:**
1. Design prediction cards for analytics dashboard
2. Implement UI components
3. Add confidence interval display
4. Add historical accuracy section
- **Deliverables:**
- Prediction UI in analytics page
- Mobile-responsive design
- **Team:** 1 developer

---

## Phase 4: Testing & Refinement (Weeks 9-10)

**Week 9: Beta Testing**
- **Tasks:**
1. Launch to 10 Dealer users (invite-only)
2. Collect feedback
3. Monitor prediction accuracy
4. Fix bugs
- **Deliverables:**
- Beta feedback report
- Bug fixes
- Accuracy monitoring dashboard
- **Success Metric:** >80% beta users find predictions valuable

**Week 10: Production Launch Prep**
- **Tasks:**
1. Finalize disclaimers and legal copy
2. Update Terms of Service
3. Create help docs and FAQs
4. Train support team
5. Set up monitoring alerts
- **Deliverables:**
- Production-ready feature
- Support documentation
- Monitoring dashboard
- **Launch:** Friday Week 10, 5pm (low-traffic time)

---

## Phase 5: Launch & Monitoring (Week 11+)

**Week 11: Public Launch**
- **Activities:**
1. Email campaign to all Dealer tier
2. Blog post + social media
3. Press release distribution
4. Monitor usage metrics
- **Success Metrics (30 days):**
- 60%+ Dealer users view predictions

- 10+ Collector → Dealer upgrades
- 85%+ prediction accuracy
- <5 support tickets/day

**Ongoing:**
- Weekly model retraining
- Monthly accuracy audits
- Quarterly data partnership reviews
- User feedback integration

---

# Risk Assessment

## High-Priority Risks

**Risk 1: Prediction Accuracy Below 80%**
- **Impact:** HIGH (damages brand credibility)
- **Probability:** MEDIUM (30%)
- **Mitigation:**
- Conservative MVP (watches only, best data)
- Large confidence intervals (±15%)
- Clear disclaimers
- Beta testing with feedback loop
- Go/no-go decision at Week 6

**Risk 2: Data Partnership Costs Exceed Budget**
- **Impact:** MEDIUM (reduces ROI)
- **Probability:** MEDIUM (40%)
- **Mitigation:**
- Negotiate upfront
- Cap at $5k/year
- Backup: Manual data collection
- Phase 2: Monetize data (sell insights to partners)

**Risk 3: Low User Adoption**
- **Impact:** HIGH (no revenue impact)
- **Probability:** LOW (20%)
- **Mitigation:**
- Strong user demand signals (18 requests)
- Natural extension of existing analytics (47% usage)
- Prominent placement in UI
- Email campaign + education

---

## Medium-Priority Risks

**Risk 4: Competitor Launches Similar Feature**
- **Impact:** MEDIUM (reduces differentiation)
- **Probability:** MEDIUM (30%)
- **Mitigation:**
- Speed to market (10 weeks)

- Patent application (ML model for luxury assets)
- Continuous improvement (stay ahead)

**Risk 5: Model Drift (Accuracy Degrades Over Time)**
- **Impact:** MEDIUM (user dissatisfaction)
- **Probability:** HIGH (60%)
- **Mitigation:**
- Automated monthly retraining
- Monitoring dashboard with alerts
- A/B testing new models before rollout

---

# Success Metrics

## Development Metrics (Weeks 1-10)

- **On-Time Delivery:** Complete all milestones within 10 weeks
- **Model Accuracy:** >80% within ±15% on test set
- **API Performance:** <500ms prediction latency
- **Bug Count:** <10 production bugs in first 30 days

## Business Metrics (First 90 Days)

- **User Engagement:** 60%+ of Dealer users view predictions weekly
- **Conversion:** 10+ Collector → Dealer upgrades citing predictions
- **Revenue:** +$2k MRR from upgrades
- **NPS Impact:** +5 points among users who use predictions

## Product Metrics (First 90 Days)

- **Adoption:** 40%+ of Dealer users view predictions (activation)
- **Retention:** 70%+ return to predictions monthly
- **Satisfaction:** 4+ star rating in user feedback
- **Support:** <5 prediction-related tickets/day

## Technical Metrics (First 90 Days)

- **Prediction Accuracy:** 85%+ maintained over time
- **API Uptime:** 99.5%+
- **Model Staleness:** Retrained within 30 days
- **Data Freshness:** Price data <7 days old

---

# Resource Allocation

## Budget Breakdown

| Category | Cost | Notes |
|---|---|---|
| Development Labor | $18,000 | 10 weeks @ $1,800/week |
| Data Scientist | $2,000 | 2-week contract for model training |
| Chrono24 API Setup | $2,000 | Initial data access |
| AWS SageMaker | $500 | First month (then $500/month ongoing) |
| Data APIs | $200 | First month (then $200/month ongoing) |
| Marketing | $5,000 | Launch campaign |
| Contingency (10%) | $2,770 | Unexpected costs |
| **TOTAL** | **$30,470** | |

**Note:** Slightly over $30k budget, but Marketing can be reduced to $2.5k if needed.

## Team Allocation

**Primary Developer:** 10 weeks full-time
- Weeks 1-2: Data engineering
- Weeks 3-6: ML development (with consultant)
- Weeks 7-8: API + UI integration
- Weeks 9-10: Testing + launch prep

**Data Scientist Consultant:** 2 weeks part-time (Weeks 4-5)
- Model architecture design
- Training supervision
- Validation methodology

**DevOps:** 5 hours (Week 3)
- AWS SageMaker setup
- Model deployment pipeline

**Support:** 2 hours training (Week 10)
- Feature overview
- FAQ review
- Escalation process

---

# Next Steps

## Immediate Actions (This Week)

1. **Approve Phase 8 Roadmap**
   - Review this document
   - Approve Predictive Analytics as Priority 1
   - Approve $30k budget

2. **Initiate Data Partnership**
   - Contact Chrono24 API sales
   - Begin contract negotiations
   - Set up demo account

3. **Technical Prep**
   - Set up AWS SageMaker account
   - Design database schema for predictions
   - Create project repo

## Week 1 Kickoff

1. **Project Kickoff Meeting**
   - Review execution plan
   - Assign tasks
   - Set up weekly check-ins

2. **Begin Development**
   - Start data ingestion pipeline
   - Chrono24 API integration
   - Database schema implementation

3. **Stakeholder Communication**
   - Email to Dealer tier: “Coming soon”
   - Internal announcement
   - Set up Slack channel: #predictive-analytics

---

# Approval Required

**Decision Maker:** CEO/Founder

**Approval Checklist:**
- [ ] Approve Predictive Analytics as Phase 8 Priority 1
- [ ] Approve $30,470 budget
- [ ] Approve 10-week timeline
- [ ] Approve risk mitigation strategies
- [ ] Approve success metrics
- [ ] Authorize data partnership negotiations
- [ ] Approve launch to Dealer tier only (initially)

**Estimated Timeline:** 10 weeks
**Estimated Cost:** $30,470
**Expected Return:** +$25-50k ARR

---

**Mission Status:** Awaiting Approval
**Next Review:** After approval, begin Week 1
**Document Version:** 1.0
**Last Updated:** December 2, 2024